

# Dawkins Weasel 1.2.0 — ODD Protocol (Overview, Design Concepts, Details)

## 1. Purpose

Dawkins Weasel is a NetLogo model designed for instructional use. It demonstrates why **cumulative selection** (retaining partial improvements) can rapidly produce a specified target phrase, while **random change alone** is extraordinarily unlikely to do so for longer phrases. The model is inspired by Richard Dawkins's "Weasel" thought experiment in *The Blind Watchmaker*.

## 2. Entities, state variables, and scales

### Entities

The model's entities are **strings** (phrases) of fixed length, represented as NetLogo strings.

- **Parent string (single entity):** the current best string, stored as `parent-string`.
- **Offspring strings (temporary entities):** strings generated each generation (either many offspring when selection is on, or an in-place mutated string when selection is off). Offspring are not stored as persistent agents; they exist as local variables during each step.

There is no spatial structure and no interaction among entities; comparison occurs only between each candidate string and the fixed target.

### Global variables

- `allowed-chars` (global): "ABCDEFGHIJKLMNOPQRSTUVWXYZ " (uppercase letters plus space).
- `parent-string` (global): current best phrase.
- `generation` (global): current generation count.

### User settings

- `target-phrase` (input): target phrase the model attempts to match.
- `mutation-rate` (slider): per-character probability of mutation during copying.
- `offspring-per-generation` (input): number of offspring generated per generation when selection is on.
- `with-selection` (switch): whether selection is applied.

### Temporal scale

Time is discrete and measured in **generations**. Each call to `go` advances the model exactly **one generation**.

### 3. Process overview and scheduling

The model proceeds in the following sequence:

#### Setup phase

1. The model clears previous state and output.
2. It defines the allowed character set (**allowed-chars**).
3. It **normalizes** **target-phrase** so it contains only characters in **allowed-chars**:
  - lowercase letters are converted to uppercase
  - any other characters are replaced with spaces
  - if **target-phrase** is empty, a default phrase is assigned
4. It creates a random initial string of the same length as **target-phrase** and stores it as **parent-string**.
5. It sets **generation = 0** and prints the initial string to the Output Area.

#### Runtime phase (one generation per tick)

On each call to **go**:

##### If **with-selection** is ON (cumulative selection):

1. Increment **generation**.
2. Generate **offspring-per-generation** offspring strings by copying **parent-string**.
3. For each offspring, iterate through all character positions:
  - with probability **mutation-rate**, replace the character with a random element of **allowed-chars**
4. Compute a score for each offspring (defined below) and select the **best** offspring (lowest score).
5. Set **parent-string** to the best offspring.
6. Print the generation and current **parent-string** to the Output Area.

##### If **with-selection** is OFF (random change only):

1. Increment **generation**.
2. Iterate through all character positions of **parent-string**:
  - with probability **mutation-rate**, replace the character with a random element of **allowed-chars**
3. Print the generation and current **parent-string** to the Output Area.

#### Stopping condition

After each generation, the model computes a score for the current string. If the score equals 0 (exact match), the model stops and prints a summary statement including:

- number of generations required

- target-phrase
- mutation-rate
- offspring-per-generation
- whether selection was on or off

## 4. Design concepts

### Basic principles

The model implements Dawkins's contrast between:

- **single-step random typing** (no retention of partial improvements), and
- **cumulative selection** (retention of partial improvements across generations).

The “fitness” concept is reduced to a single operational measure: similarity to a fixed target phrase.

### Emergence

When selection is on, the model typically produces rapid convergence to the target phrase under intermediate parameter settings. Key emergent behaviors include:

- **No progress when mutation-rate = 0:** no variation is produced; selection has nothing to act on.
- **Slowed or unstable progress when mutation-rate is too high:** improvements are frequently overwritten.
- **Faster convergence with larger offspring-per-generation (when selection is on):** more candidate variants increase the chance of producing a better match in each generation.
- **Near-impossibility of convergence when selection is off for long phrases:** the expected time to match grows explosively with phrase length.

### Adaptation

Adaptation is represented as a reduction in the number of mismatched characters between the current string and the target phrase over generations, due to selection retaining higher-matching offspring.

### Objectives

The objective function is to minimize mismatch to the target phrase, computed as: **score = number of positions where parent-string differs from target-phrase**

### Learning

No learning occurs. Strings change only through mutation and (optionally) selection.

## Prediction

No prediction in the forecasting sense is performed; however, the model produces expected qualitative outcomes:

- With selection on and moderate `mutation-rate`, convergence typically occurs in relatively few generations.
- Without selection, convergence is extremely unlikely for longer phrases.

## Sensing

Strings do not “sense” in an agent-based sense. The model algorithm computes similarity to the target phrase and uses it for selection.

## Interaction

There is no interaction among strings; each candidate is compared only to the target phrase.

## Stochasticity

Randomness enters in two places:

- whether each character mutates (`random-float 1.0 < mutation-rate`)
- the replacement character chosen from `allowed-chars`

## Observation

The Output Area reports one line per generation: generation number and the current best string. A final summary statement is printed in the Command Center upon reaching an exact match.

## 5. Initialization

At setup:

- `allowed-chars` is set to "`ABCDEFGHIJKLMNOPQRSTUVWXYZ`".
- `target-phrase` is normalized to use only `allowed-chars` and to be non-empty.
- `parent-string` is initialized as a random string of the same length as `target-phrase`, with each character sampled uniformly from `allowed-chars`.
- `generation` starts at 0.

## 6. Input data

No external input data files are used. Inputs are provided through the NetLogo interface (`target-phrase`, `mutation-rate`, `offspring-per-generation`, `with-selection`).

## 7. Submodels

### Mutation submodel

For each character position, with probability `mutation-rate`, the character is replaced by a uniformly random character from `allowed-chars`.

### Selection submodel (conditional on `with-selection`)

Among `offspring-per-generation` mutated offspring, the model selects the offspring with the lowest mismatch score relative to `target-phrase`. (If ties occur, the first encountered tied best is retained under the current implementation.)

### Scoring submodel

Similarity is computed position-by-position across the whole string:

- each exact character match reduces the mismatch count by 1
- the final score equals the number of mismatches